

Section 4 Grouping Data

Introduction to PostgreSQL GROUP BY clause

The GROUP BY clause divides the rows returned from the [SELECT](#) statement into groups. For each group, you can apply an aggregate function e.g., [SUM\(\)](#) to calculate the sum of items or [COUNT\(\)](#) to get the number of items in the groups.

The following statement illustrates the basic syntax of the GROUP BY clause:

```
SELECT
    column_1,
    column_2,
    ...,
    aggregate_function(column_3)
FROM
    table_name
GROUP BY
    column_1,
    column_2,
    ...;
```

In this syntax:

- First, select the columns that you want to group e.g., column1 and column2, and column that you want to apply an aggregate function (column3).
- Second, list the columns that you want to group in the GROUP BY clause.

The statement clause divides the rows by the values of the columns specified in the GROUP BY clause and calculates a value for each group.

It's possible to use other clauses of the SELECT statement with the GROUP BY clause.

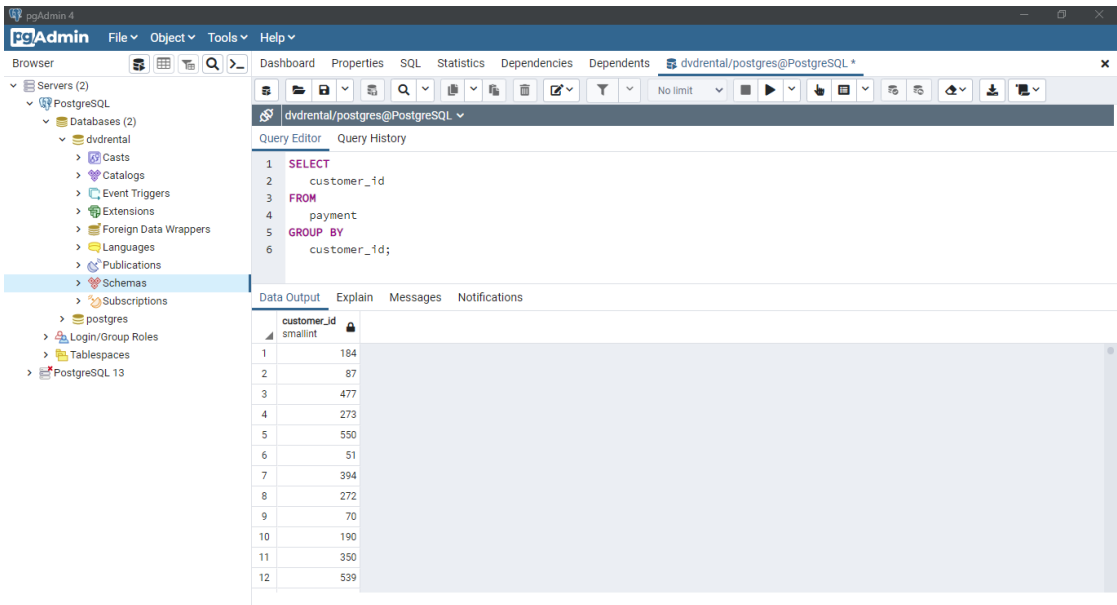


PostgreSQL evaluates the GROUP BY clause after the FROM and [WHERE](#) clauses and before the [HAVING](#) [SELECT](#), [DISTINCT](#), ORDER BY and [LIMIT](#) clauses.

PostgreSQL GROUP BY clause

1) Using PostgreSQL GROUP BY without an aggregate function example

You can use the GROUP BY clause without applying an aggregate function. The following query gets data from the payment table and groups the result by customer id.



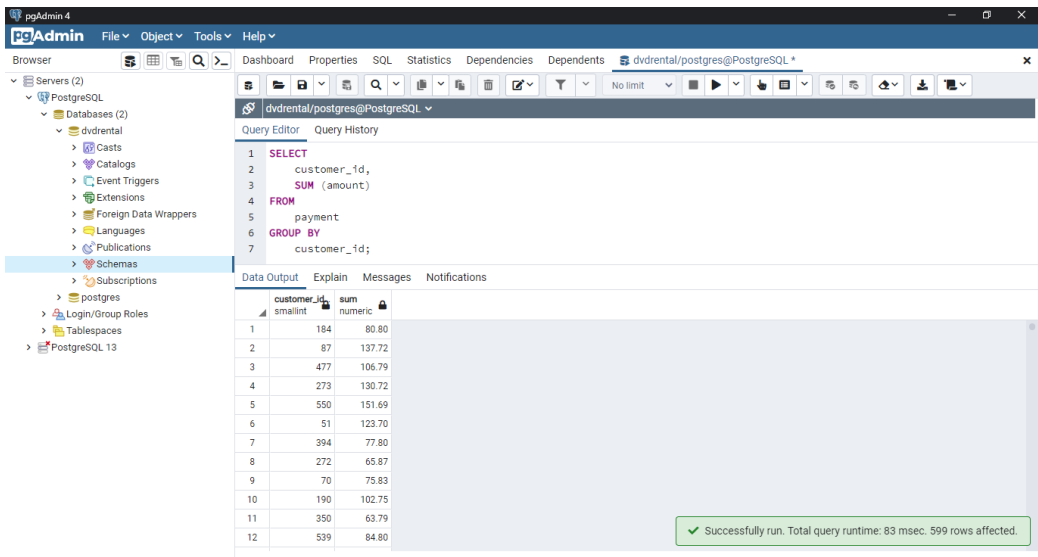
In this case, the GROUP BY works like the DISTINCT clause that removes duplicate rows from the result set.

2) Using PostgreSQL GROUP BY with SUM () function example

The GROUP BY clause is useful when it is used in conjunction with an [aggregate function](#).

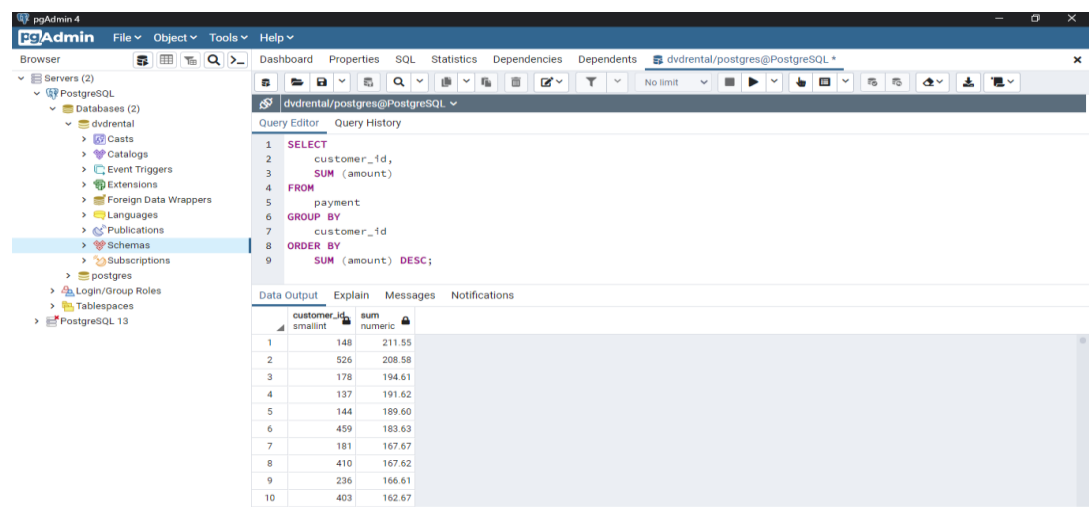
For example, to select the total amount that each customer has been paid, you use the GROUP BY clause to divide the rows in the payment table into groups grouped by customer id. For each group, you calculate the total amounts using the [SUM \(\)](#) function.

The following query uses the GROUP BY clause to get total amount that each customer has been paid:



The GROUP BY clause sorts the result set by customer id and adds up the amount that belongs to the same customer. Whenever the customer_id changes, it adds the row to the returned result set.

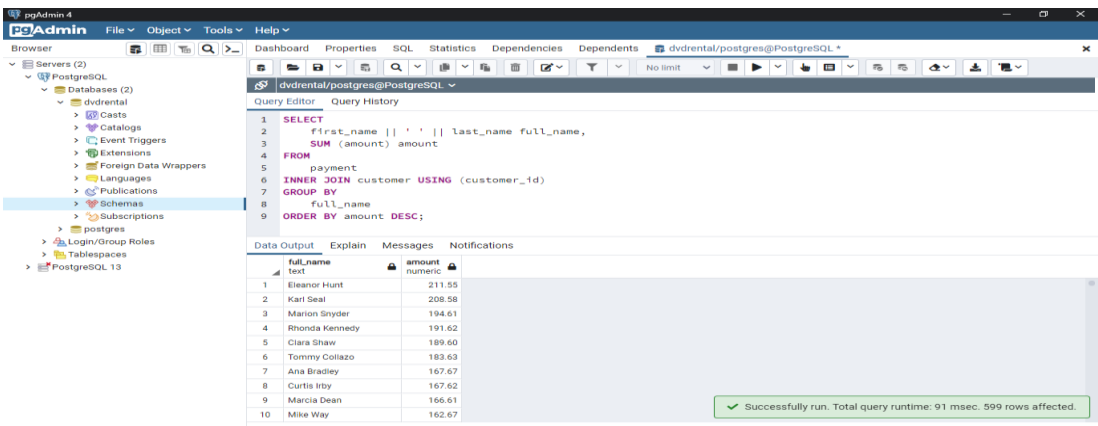
The following statement uses the ORDER BY clause with GROUP BY clause to sort the groups:



3) Using PostgreSQL GROUP BY clause with the JOIN clause

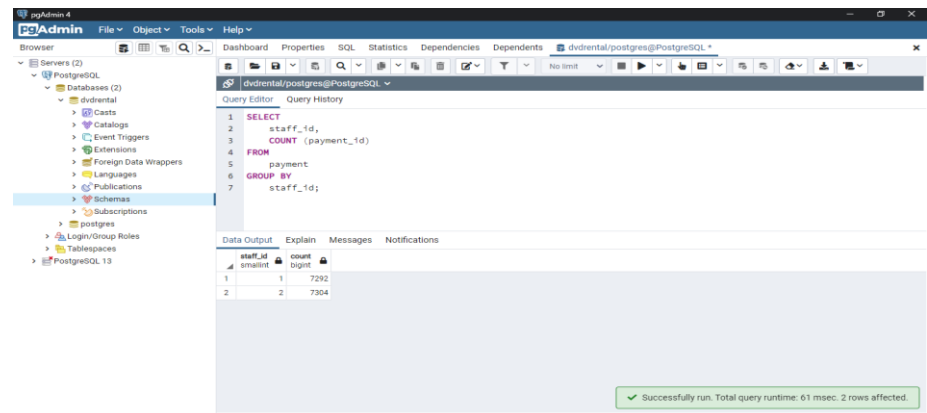
The following statement uses the GROUP BY clause with the INNER JOIN clause the get the total amount paid by each customer.

Unlike the previous example, this query joins the payment table with the customer table and group customers by their names.



4) Using PostgreSQL GROUP BY with COUNT () function example

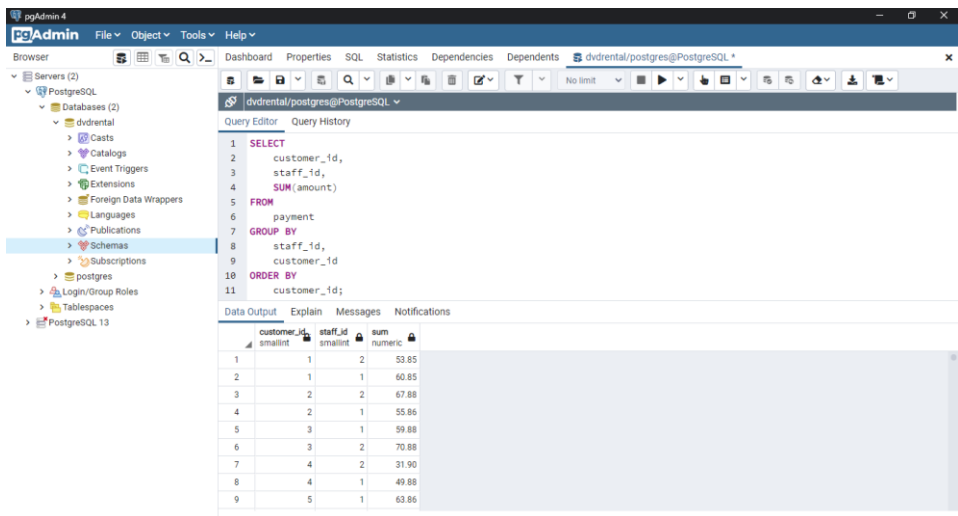
To find the number of payment transactions that each staff has been processed, you group the rows in the payment table by the values in the staff_id column and use the COUNT () function to get the number of transactions:



The GROUP BY clause divides the rows in the payment into groups and groups them by value in the staff_id column. For each group, it returns the number of rows by using the COUNT () function.

5) Using PostgreSQL GROUP BY with multiple columns

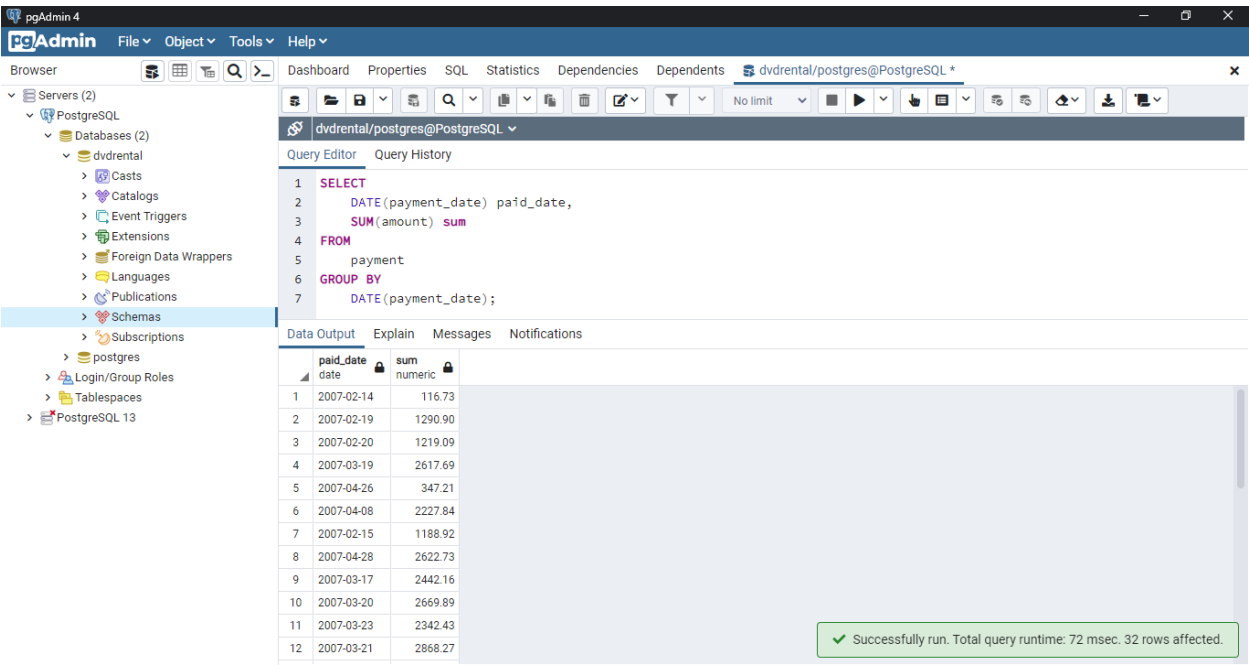
The following example uses multiple columns in the GROUP BY clause:



In this example, the GROUP BY clause divides the rows in the payment table by the values in the customer_id and staff_id columns. For each group of (customer_id, staff_id), the SUM () calculates the total amount.

6) Using PostgreSQL GROUP BY clause with date column

The payment_date is a timestamp column. To group payments by dates, you use the DATE () function to convert timestamps to dates first and then group payments by the result date:



Introduction to PostgreSQL HAVING clause

The HAVING clause specifies a search condition for a group or an aggregate. The HAVING clause is often used with the GROUP BY clause to filter groups or aggregates based on a specified condition.

The following statement illustrates the basic syntax of the HAVING clause:

```
SELECT
    column1,
    aggregate_function (column2)
FROM
    table_name
GROUP BY
    column1
HAVING
    condition;
```

In this syntax, the group by clause returns rows grouped by the column1. The HAVING clause specifies a condition to filter the groups.

It's possible to add other clauses of the SELECT statement such as JOIN, LIMIT, and FETCH etc.

PostgreSQL evaluates the HAVING clause after the FROM, WHERE, GROUP BY, and before the SELECT, DISTINCT, ORDER BY and LIMIT clauses.



Since the HAVING clause is evaluated before the SELECT clause, you cannot use column aliases in the HAVING clause. Because at the time of evaluating the HAVING clause, the column aliases specified in the SELECT clause are not available.

HAVING vs. WHERE

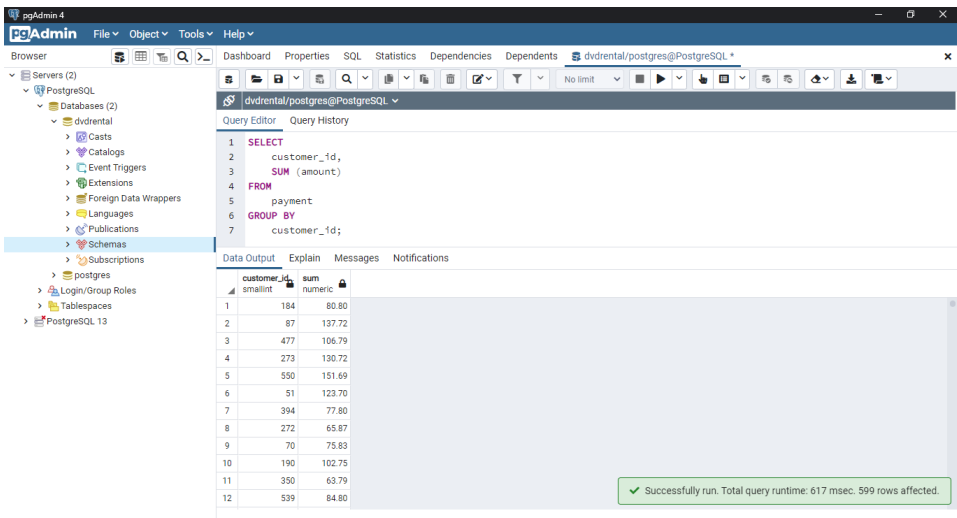
The WHERE clause allows you to filter rows based on a specified condition. However, the HAVING clause allows you to filter groups of rows according to a specified condition.

In other words, the WHERE clause is applied to rows while the HAVING clause is applied to groups of rows.

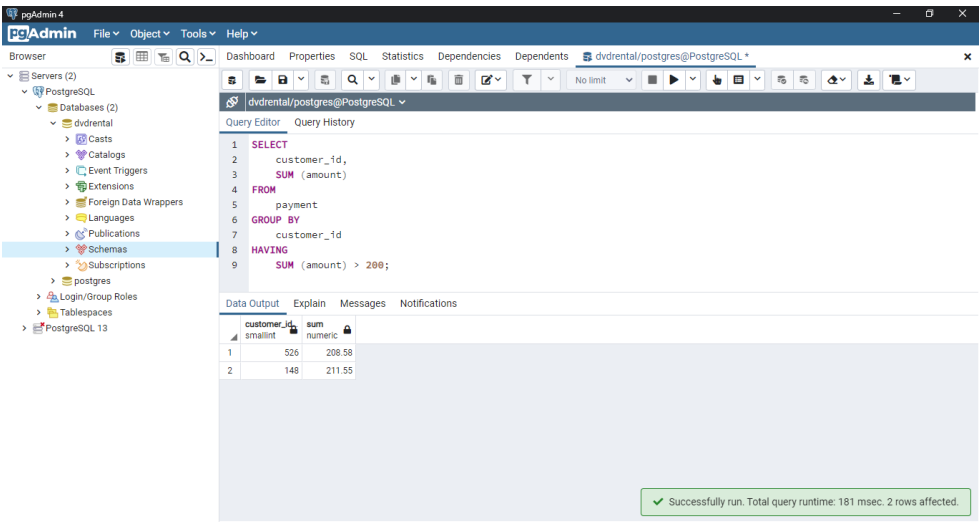
PostgreSQL HAVING clause examples

1) Using PostgreSQL HAVING clause with SUM function example

The following query uses the GROUP BY clause with the SUM () function to find the total amount of each customer:

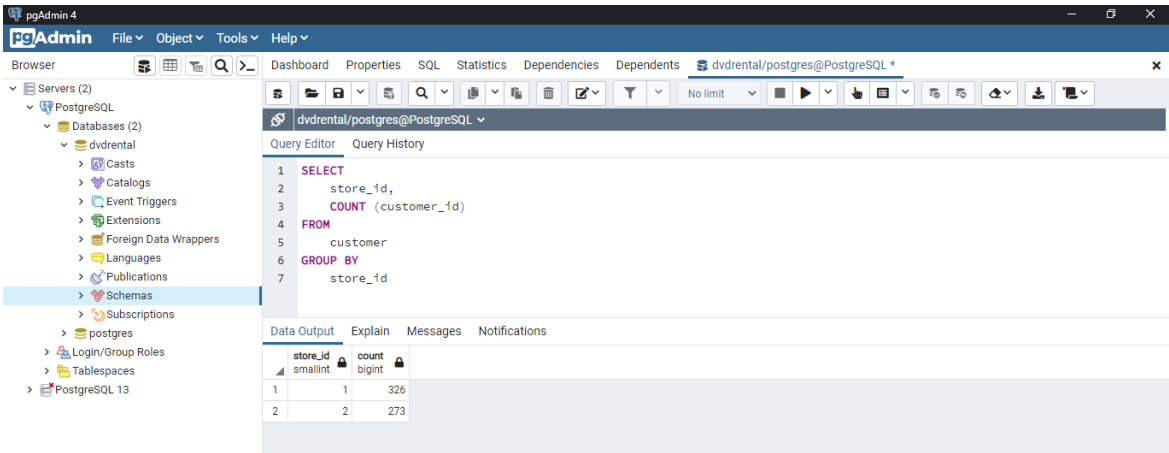


The following statement adds the HAVING clause to select the only customers who have been spending more than 200:



2) PostgreSQL HAVING clause with COUNT example

The following query uses the GROUP BY clause to find the number of customers per store:



The following statement adds the HAVING clause to select the store that has more than 300 customers:

pgAdmin 4

pgAdmin

FileObjectToolsHelp

Browser

Servers (2)

PostgreSQL

Databases (2)

dvdrental

Casts

Catalogs

Event Triggers

Extensions

Foreign Data Wrappers

Languages

Publications

Schemas

Subscriptions

postgres

Login/Group Roles

Tablespaces

PostgreSQL 13

DashboardPropertiesSQLStatisticsDependenciesDependents

dvdrental/postgres@PostgreSQL *

Query EditorQuery History

1

2

3

4

5

6

7

8

9

SELECT

store_id,

COUNT (customer_id)

FROM

customer

GROUP BY

store_id

HAVING

COUNT (customer_id) > 300;

Data Output

Explain

Messages

Notifications

store_id

count

1

326

store_id	count
1	326