

## Introduction to PostgreSQL UNION operator

The UNION operator combines result sets of two or more SELECT statements into a single result set.

The following illustrates the syntax of the UNION operator that combines result sets from two queries:

```
SELECT select_list_1
FROM table_expression_1
UNION
SELECT select_list_2
FROM table_expression_2
```

To combine the result sets of two queries using the UNION operator, the queries must conform to the following rules:

The number and the order of the columns in the select list of both queries must be the same.

The data types must be compatible.

The UNION operator removes all duplicate rows from the combined data set. To retain the duplicate rows, you use the UNION ALL instead.

The following Venn diagram illustrates how the UNION works:



### PostgreSQL UNION with ORDER BY clause

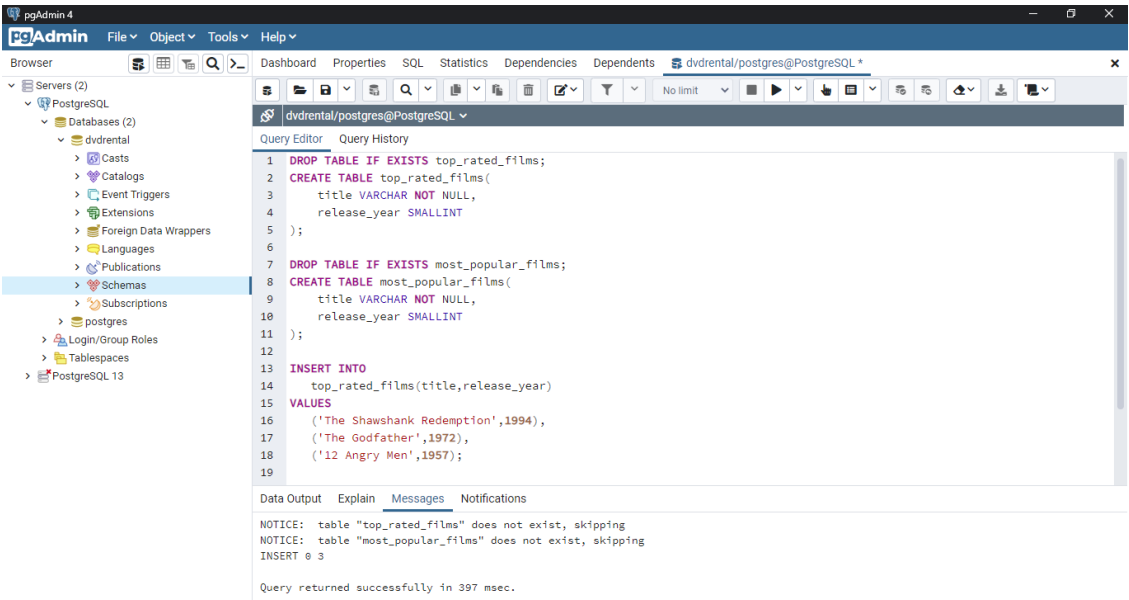
The UNION operator may place the rows from the result set of the first query before, after, or between the rows from the result set of the second query.

To sort rows in the final result set, you use the ORDER BY clause in the second query.

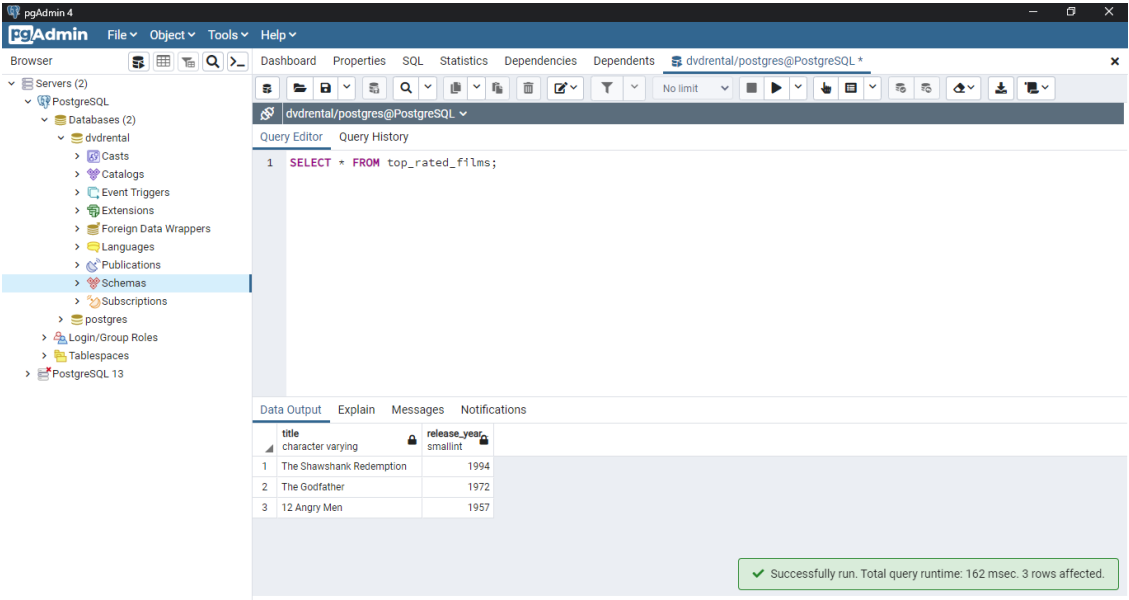
In practice, you often use the UNION operator to combine data from similar tables, which are not perfectly normalized, in the data warehouse or business intelligence systems.

### Setting up sample tables

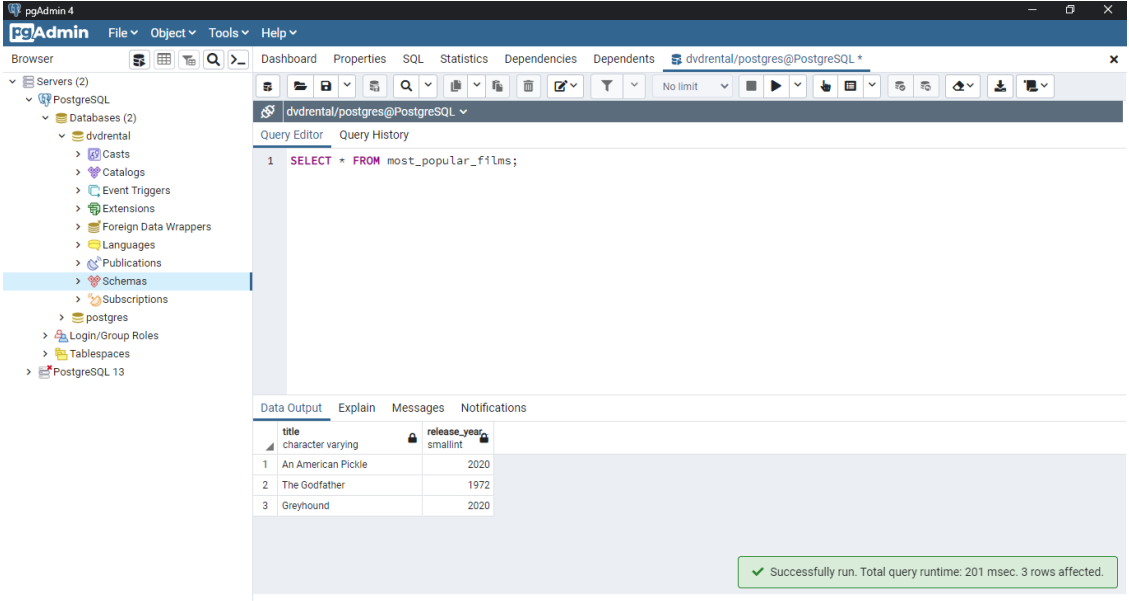
The following statements create two tables: `topRatedFilms` and `mostPopularFilms`, and insert data into these tables:



The following shows the data from the top\_rated\_films table:



The following statement returns the data from the most\_popular\_films table:

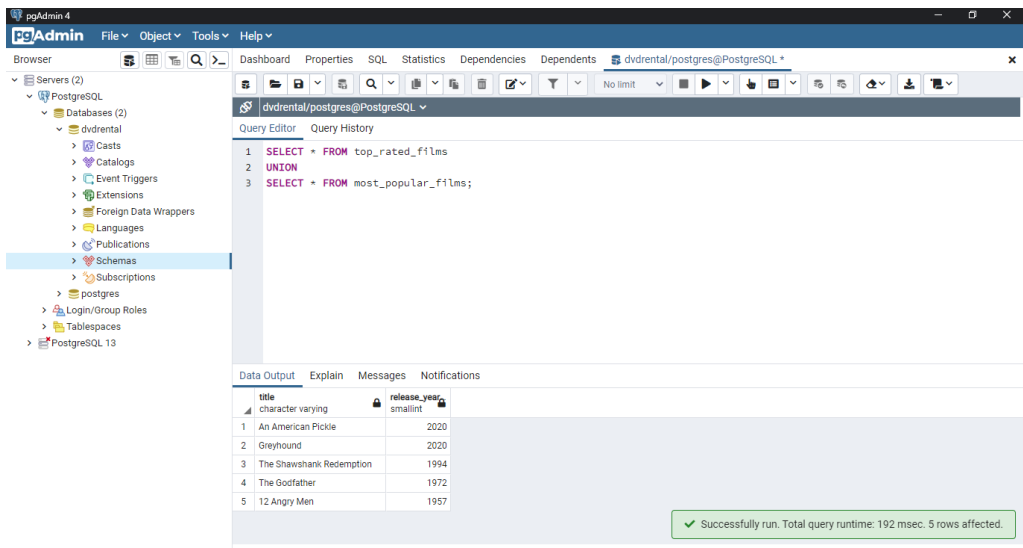


## PostgreSQL UNION examples

Let’s take some examples of using the PostgreSQL UNION operator.

### 1) Simple PostgreSQL UNION example

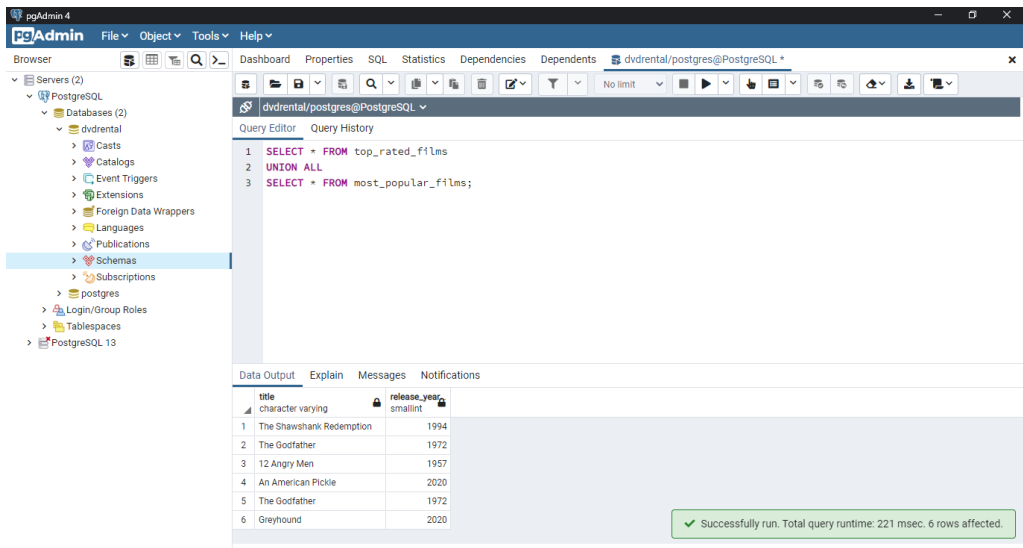
The following statement uses the UNION operator to combine data from both tables:



The result set includes five rows in the result set because the UNION operator removes one duplicate row.

### 2) PostgreSQL UNION ALL example

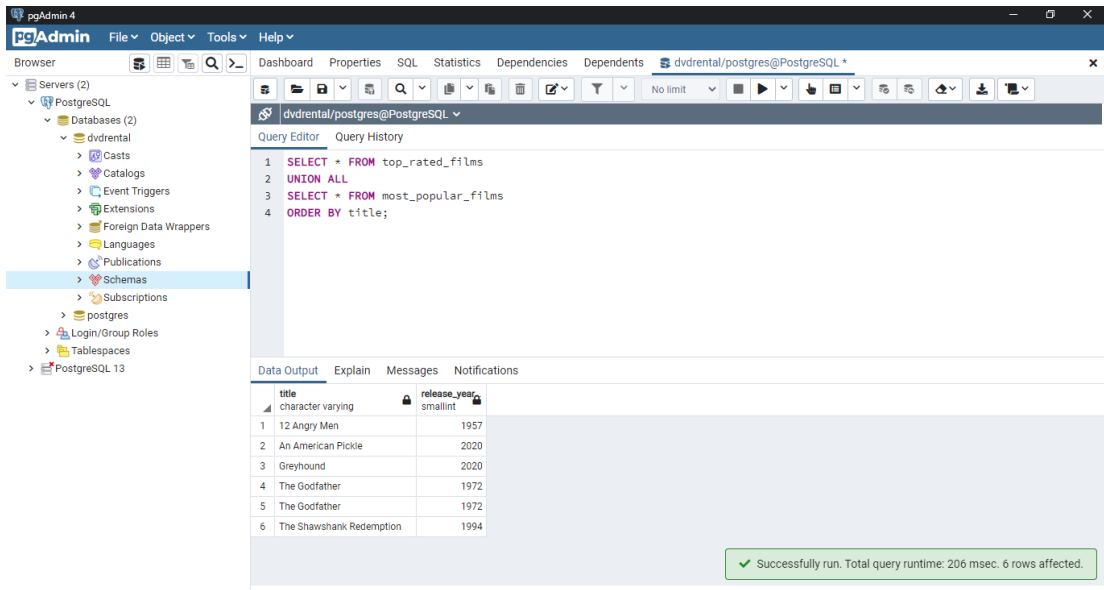
The following statement uses the UNION ALL operator to combine result sets from the top-rated\_films and most-popular\_films tables:



In this example, the duplicate row is retained in the result set.

### 3) PostgreSQL UNION ALL with ORDER BY clause example

To sort the result returned by the UNION operator, you place the ORDER BY clause end of the last query like this:



If you place the ORDER BY clause at the end of each query, the combined result set will not be sorted as you expected.


Because when UNION operator combines the sorted result sets from each query, it does not guarantee the order of rows in the final result set.

**Introduction to PostgreSQL INTERSECT operator**

Like the UNION and EXCEPT operators, the PostgreSQL INTERSECT operator combines result sets of two or more SELECT statements into a single result set.

The INTERSECT operator returns any rows that are available in both result sets.

The following illustration shows the final result set produced by the INTERSECT operator.



The final result set is represented by the yellow area where circle A intersects with circle B.

The following illustrates the syntax of the INTERSECT operator:

```
SELECT select_list
FROM A
INTERSECT
SELECT select_list
FROM B;
```

To use the INTERSECT operator, the columns that appear in the SELECT statements must follow the following rules:

The number of columns and their order in the SELECT clauses must be the same.

The data types of the columns must be compatible.

**PostgreSQL INTERSECT with ORDER BY clause**

If you want to sort the result set returned by the INTERSECT operator, you place the ORDER BY at the final query in the query list like this:

```
SELECT select_list
FROM A
INTERSECT
SELECT select_list
FROM B
ORDER BY sort_expression;
```

**PostgreSQL INTERSECT operator examples**

We'll use the top Rated films and most popular films tables created in the UNION tutorial:

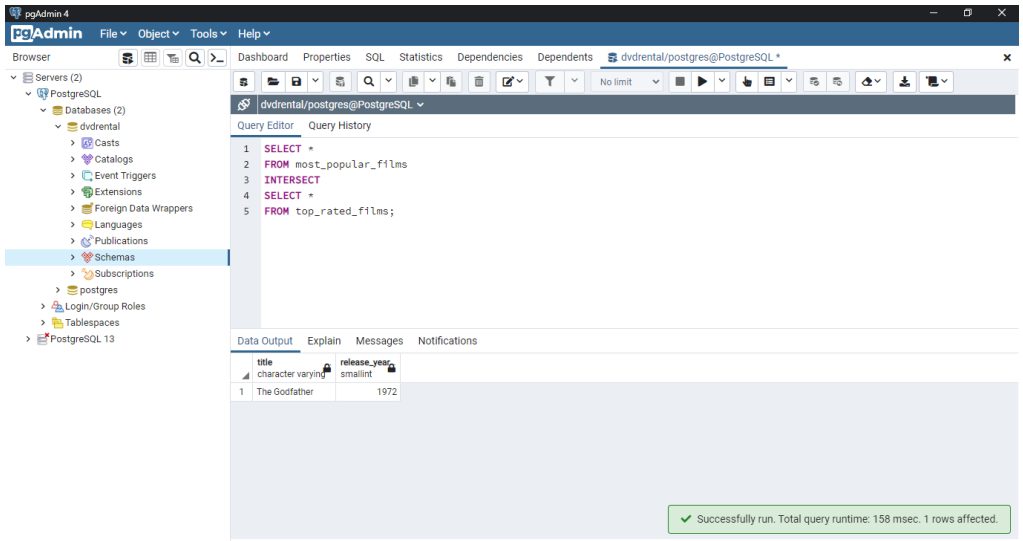
The top Rated films table:

|   | title<br>character varying | release_year<br>smallint |
|---|----------------------------|--------------------------|
| 1 | The Shawshank Redemption   | 1994                     |
| 2 | The Godfather              | 1972                     |
| 3 | 12 Angry Men               | 1957                     |

The most popular films table:

|   | title<br>character varying | release_year<br>smallint |
|---|----------------------------|--------------------------|
| 1 | An American Pickle         | 2020                     |
| 2 | The Godfather              | 1972                     |
| 3 | Greyhound                  | 2020                     |

To get popular films which are also top rated films, you use the INTERSECT operator as follows:



The result set returns one film that appears on both tables.

**Introduction to the PostgreSQL EXCEPT operator**

Like the UNION and INTERSECT operators, the EXCEPT operator returns rows by comparing the result sets of two or more queries.

The EXCEPT operator returns distinct rows from the first (left) query that are not in the output of the second (right) query.

The following illustrates the syntax of the EXCEPT operator.


```
SELECT select_list
FROM A
EXCEPT
SELECT select_list
FROM B;
```

The queries that involve in the EXCEPT need to follow these rules:

The number of columns and their orders must be the same in the two queries.

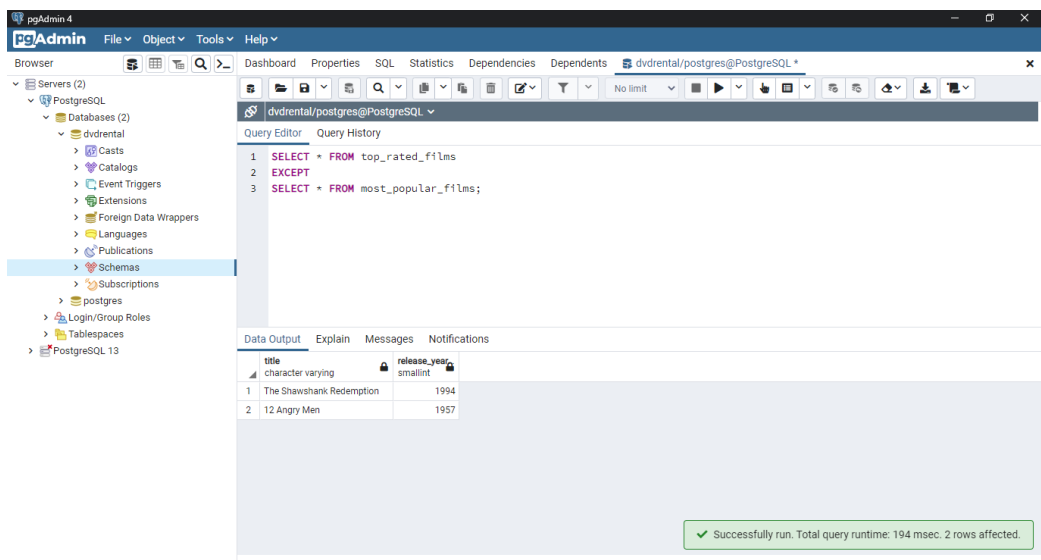
The data types of the respective columns must be compatible.

The following Venn diagram illustrates the EXCEPT operator:

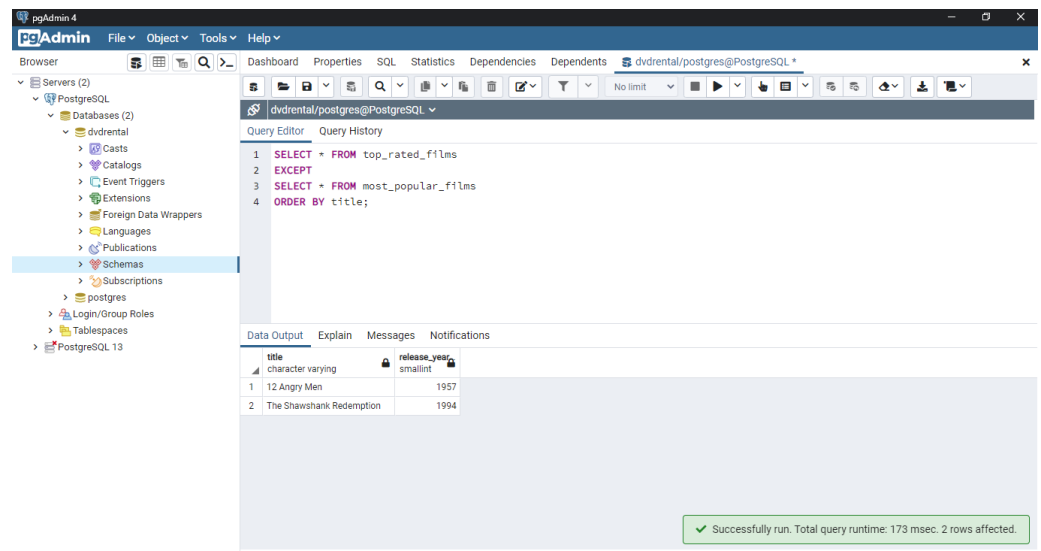


## PostgreSQL EXCEPT operator examples

The following statement uses the EXCEPT operator to find the top-rated films that are not popular:



The following statement uses the ORDER BY clause in the query to sort result sets returned by the EXCEPT operator:



Notice that we placed the ORDER BY clause at the end of the statement to sort films by title.