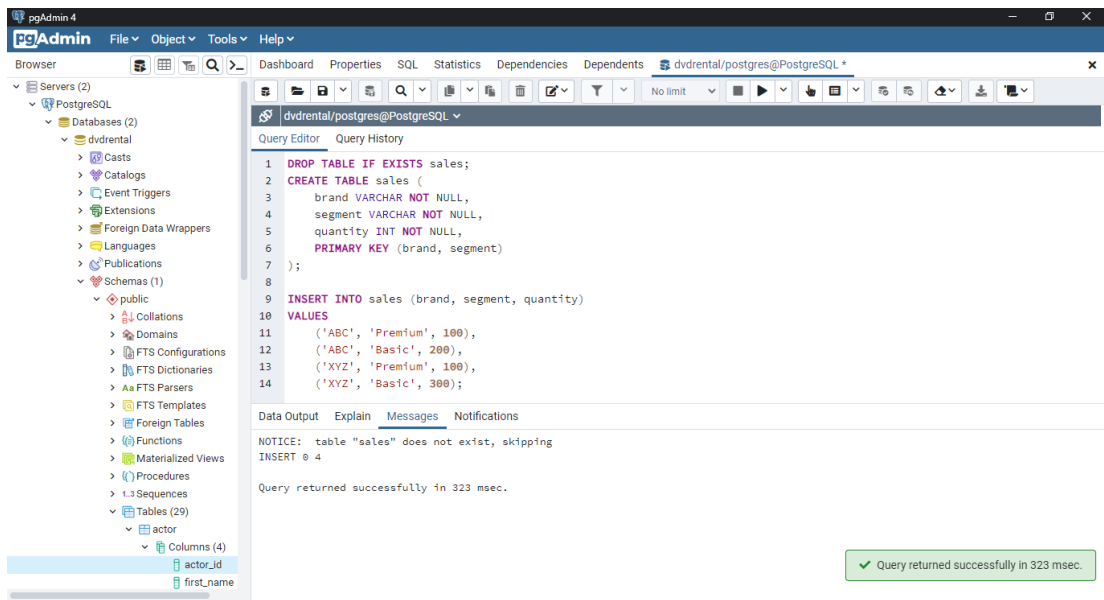


PostgreSQL GROUPING SETS

Setup a sample table

Let’s get started by creating a new table called sales for the demonstration.



The sales table stores the number of products sold by brand and segment.

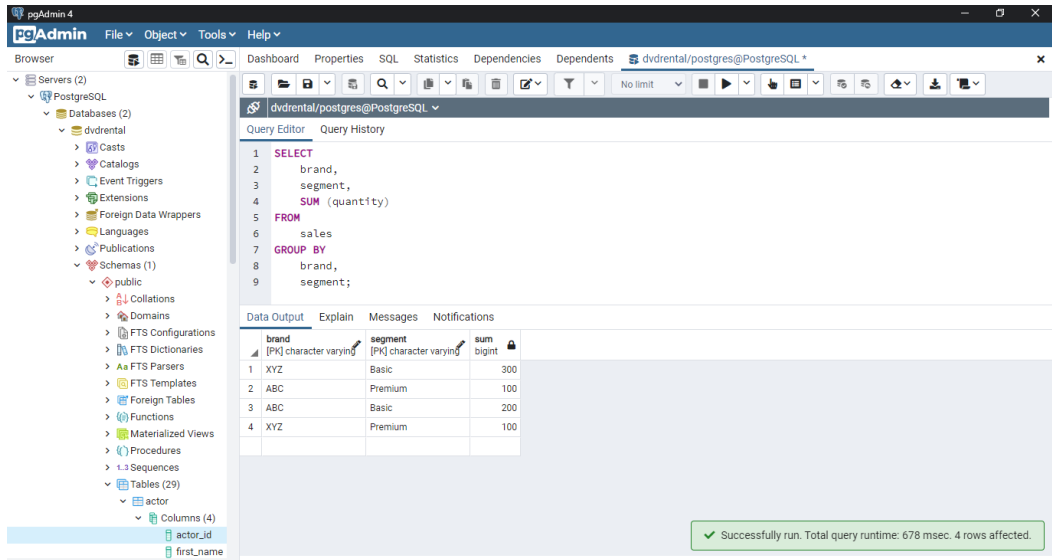
Introduction to PostgreSQL GROUPING SETS

A grouping set is a set of columns by which you group by using the GROUP BY clause.

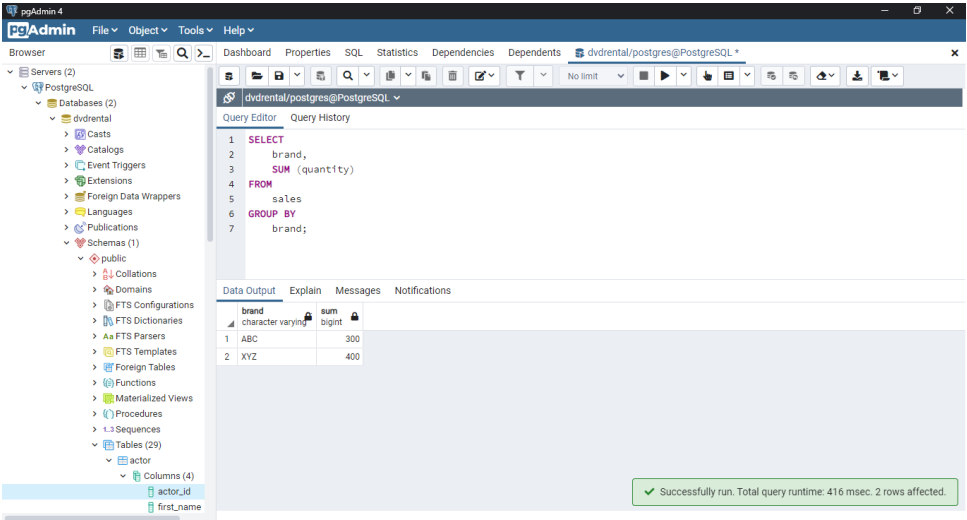
A grouping set is denoted by a comma-separated list of columns placed inside parentheses:

```
(column1, column2, ...)
```

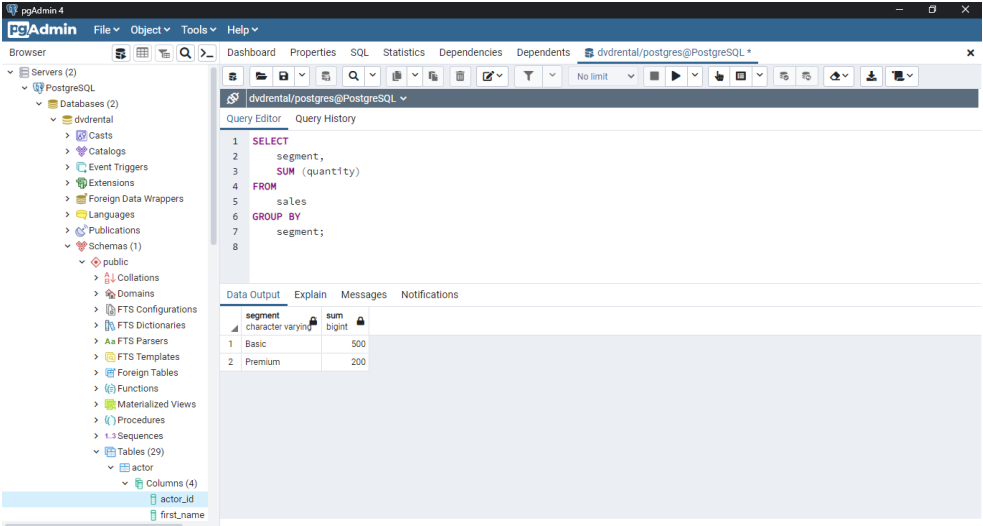
For example, the following query uses the GROUP BY clause to return the number of products sold by brand and segment. In other words, it defines a grouping set of the brand and segment which is denoted by (brand, segment)



The following query finds the number of products sold by a brand. It defines a grouping set (brand):



The following query finds the number of products sold by segment. It defines a grouping set (segment):

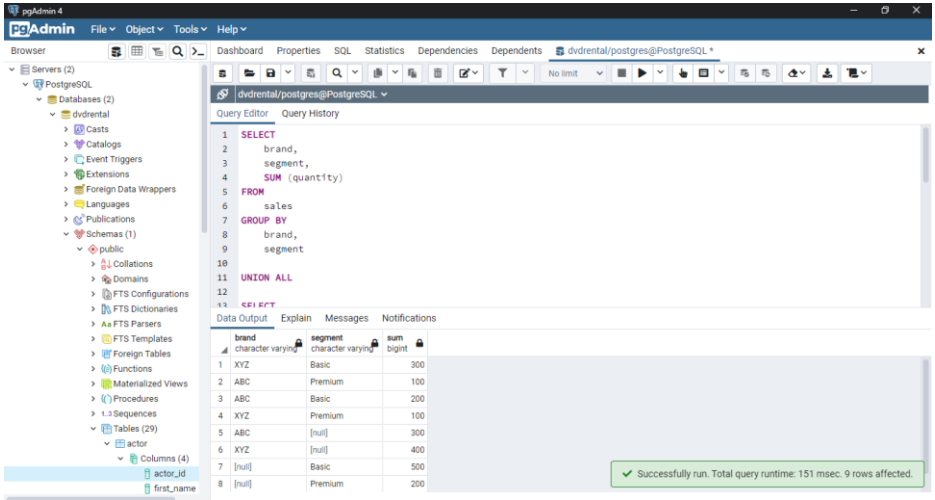


The following query finds the number of products sold for all brands and segments. It defines an empty grouping set which is denoted by ().

```
SELECT SUM (quantity) FROM sales;
```

Suppose that you want to all the grouping sets by using a single query. To achieve this, you may use the UNION ALL to combine all the queries above.

Because UNION ALL requires all result sets to have the same number of columns with compatible data types, you need to adjust the queries by adding NULL to the selection list of each as shown below:



This query generated a single result set with the aggregates for all grouping sets.

Even though the above query works as you expected, it has two main problems.

First, it is quite lengthy.

Second, it has a performance issue because PostgreSQL has to scan the sales table separately for each query.

To make it more efficient, PostgreSQL provides the GROUPING SETS clause which is the sub clause of the GROUP BY clause.

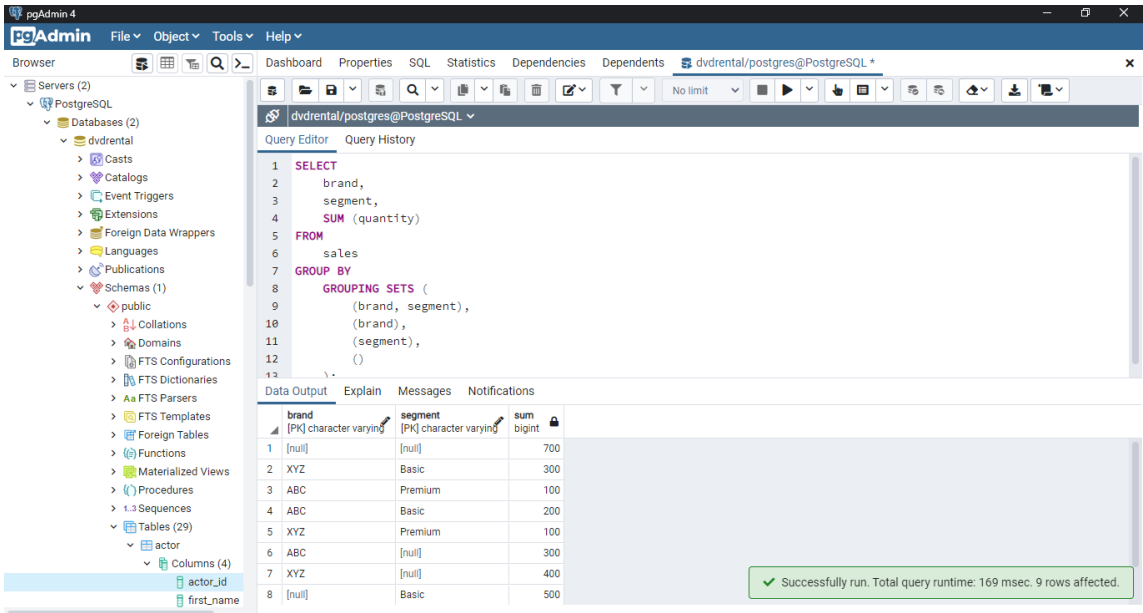
The GROUPING SETS allows you to define multiple grouping sets in the same query.

The general syntax of the GROUPING SETS is as follows:

```
SELECT
    c1,
    c2,
    aggregate_function(c3)
FROM
    table_name
GROUP BY
    GROUPING SETS (
        (c1, c2),
        (c1),
        (c2),
        ()
    );
```

In this syntax, we have four grouping sets (c1, c2), (c1), (c2), and ().

To apply this syntax to the above example, you can use GROUPING SETS clause instead of the UNION ALL clause like this:



This query is much shorter and more readable. In addition, PostgreSQL will optimize the number of times it scans the sales table and will not scan multiple times.

Grouping function

The GROUPING () function accepts an argument which can be a column name or an expression:

```
GROUPING( column_name | expression)
```

The column_name or expression must match with the one specified in the GROUP BY clause.

The GROUPING () function returns bit 0 if the argument is a member of the current grouping set and 1 otherwise.

See the following example:

The screenshot shows the pgAdmin 4 interface with a SQL query editor. The query is as follows:

```
1 SELECT
2   GROUPING(brand) grouping_brand,
3   GROUPING(segment) grouping_segment,
4   brand,
5   segment,
6   SUM (quantity)
7 FROM
8   sales
9 GROUP BY
10  GROUPING SETS (
11    (brand),
12    (segment),
13    ()
14  )
```

The 'Data Output' tab shows the following results:

grouping_brand	grouping_segment	brand	segment	sum
0	0	ABC	[null]	300
0	0	XYZ	[null]	400
1	0	[null]	Basic	500
1	0	[null]	Premium	200
1	1	[null]	[null]	700

A status message at the bottom indicates: "Successfully run. Total query runtime: 126 msec. 5 rows affected."

As shown in the screenshot, when the value in the grouping_brand is 0, the sum column shows the subtotal of the brand.

When the value in the grouping_segment is zero, the sum column shows the subtotal of the segment.

You can use the GROUPING () function in the HAVING clause to find the subtotal of each brand like this:

The screenshot shows the pgAdmin 4 interface with a SQL query editor. The query is as follows:

```
1 SUM (quantity)
2 FROM
3   sales
4 GROUP BY
5   GROUPING SETS (
6     (brand),
7     (segment),
8     ()
9   )
10 HAVING GROUPING(brand) = 0
11 ORDER BY
12   brand,
13   segment;
```

The 'Data Output' tab shows the following results:

grouping_brand	grouping_segment	brand	segment	sum
0	0	ABC	[null]	300
0	0	XYZ	[null]	400

A status message at the bottom indicates: "Successfully run. Total query runtime: 291 msec. 2 rows affected."

PostgreSQL CUBE

Introduction to the PostgreSQL CUBE

PostgreSQL CUBE is a subclause of the GROUP BY clause. The CUBE allows you to generate multiple grouping sets.

A grouping set is a set of columns to which you want to group. For more information on the grouping sets, check it out the GROUPING SETS tutorial.

The following illustrates the syntax of the CUBE subclause:

```
SELECT
    c1,
    c2,
    c3,
    aggregate (c4)
FROM
    table_name
GROUP BY
    CUBE (c1, c2, c3);
```

In this syntax:

- First, specify the CUBE subclause in the the [GROUP BY](#) clause of the SELECT statement.
- Second, in the select list, specify the columns (dimensions or dimension columns) which you want to analyze and [aggregation function](#) expressions.
- Third, in the GROUP BY clause, specify the dimension columns within the parentheses of the CUBE subclause.

The query generates all possible grouping sets based on the dimension columns specified in CUBE. The CUBE subclause is a short way to define multiple grouping sets so the following are equivalent:

```
CUBE(c1,c2,c3)

GROUPING SETS (
    (c1,c2,c3),
    (c1,c2),
    (c1,c3),
    (c2,c3),
    (c1),
    (c2),
    (c3),
    ()
)
```

In general, if the number of columns specified in the CUBE is n , then you will have 2^n combinations.

PostgreSQL allows you to perform a partial cube to reduce the number of aggregates calculated. The following shows the syntax:

```
SELECT
    c1,
    c2,
    c3,
    aggregate (c4)
FROM
    table_name
GROUP BY
    c1,
    CUBE (c1, c2);
```

PostgreSQL CUBE examples

We will use the sales table created in the GROUPING SETS tutorial for the demonstration.

brand	segment	quantity
ABC	Premium	100
ABC	Basic	200
XYZ	Premium	100
XYZ	Basic	300

The following query uses the CUBE subclause to generate multiple grouping sets:

The screenshot displays the pgAdmin 4 web interface. On the left, the 'Browser' pane shows a tree structure of the database: Servers (2) > PostgreSQL > Databases (2) > dvdrental. Under 'Schemas (1)', the 'public' schema is expanded, showing various database objects like Collations, Domains, FTS Configurations, etc. The 'Tables (29)' folder is also expanded, showing the 'actor' table.

The main pane is titled 'dvdrental/postgres@PostgreSQL *'. It contains a 'Query Editor' with the following SQL query:

```

1  SELECT
2      brand,
3      segment,
4      SUM (quantity)
5  FROM
6      sales
7  GROUP BY
8      CUBE (brand, segment)
9  ORDER BY
10     brand,
11     segment;
12

```

Below the query editor, the 'Data Output' tab is active, showing the results of the query. The results are displayed in a table with four columns: brand, segment, sum, and bigint. The data is as follows:

	brand [PK] character varying	segment [PK] character varying	sum bigint
1	ABC	Basic	200
2	ABC	Premium	100
3	ABC	[null]	300
4	XYZ	Basic	300
5	XYZ	Premium	100
6	XYZ	[null]	400
7	[null]	Basic	500
8	[null]	Premium	200

At the bottom right, a green status bar indicates: 'Successfully run. Total query runtime: 205 msec. 9 rows affected.'

PostgreSQL ROLLUP

Introduction to the PostgreSQL ROLLUP

The PostgreSQL ROLLUP is a subclause of the GROUP BY clause that offers a shorthand for defining multiple grouping sets. A grouping set is a set of columns by which you group. Check out the grouping sets tutorial for the detailed information.

Different from the CUBE subclause, ROLLUP does not generate all possible grouping sets based on the specified columns. It just makes a subset of those.

The ROLLUP assumes a hierarchy among the input columns and generates all grouping sets that make sense considering the hierarchy. This is the reason why ROLLUP is often used to generate the subtotals and the grand total for reports.

For example, the CUBE (c1, c2, c3) makes all eight possible grouping sets:

```
(c1, c2, c3)
(c1, c2)
(c2, c3)
(c1, c3)
(c1)
(c2)
(c3)
()
```

However, the ROLLUP (c1, c2, c3) generates only four grouping sets, assuming the hierarchy $c1 > c2 > c3$ as follows:

```
(c1, c2, c3)
(c1, c2)
(c1)
()
```

A common use of ROLLUP is to calculate the aggregations of data by year, month, and date, considering the hierarchy $\text{year} > \text{month} > \text{date}$

The following illustrates the syntax of the PostgreSQL ROLLUP:

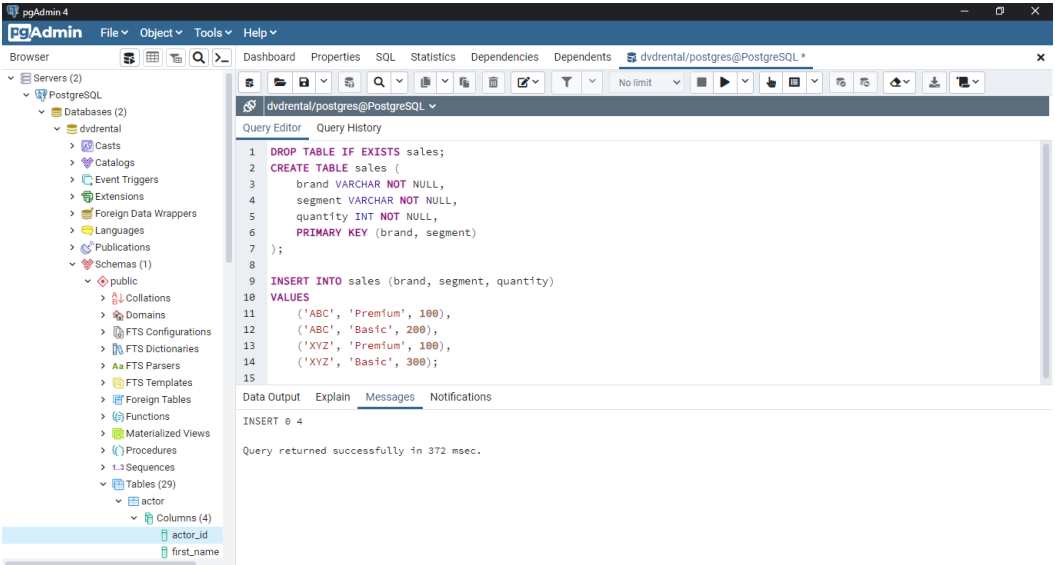
```
SELECT
    c1,
    c2,
    c3,
    aggregate(c4)
FROM
    table_name
GROUP BY
    ROLLUP (c1, c2, c3);
```

It is also possible to do a partial roll up to reduce the number of subtotals generated.

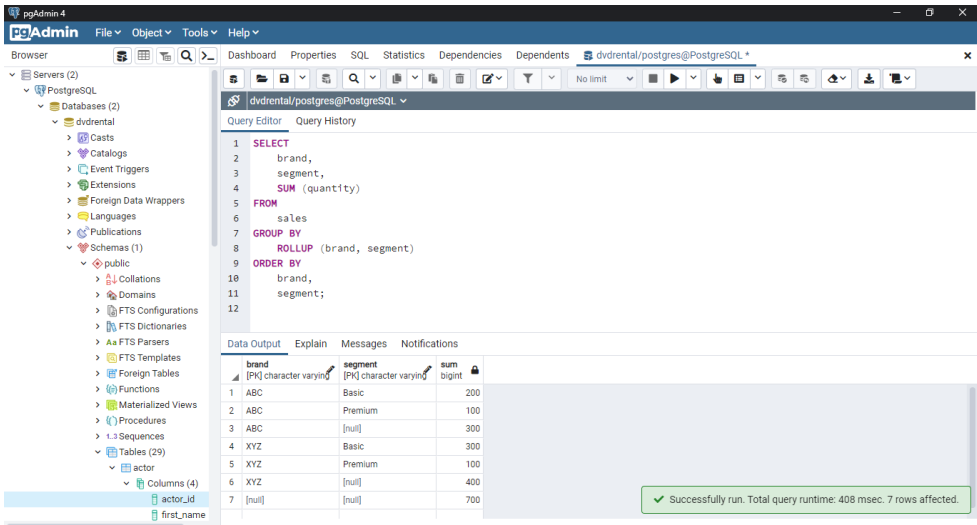
```
SELECT
    c1,
    c2,
    c3,
    aggregate(c4)
FROM
    table_name
GROUP BY
    c1,
    ROLLUP (c2, c3);
```

PostgreSQL ROLLUP examples

If you haven’t create the sales table, you can use the following script:

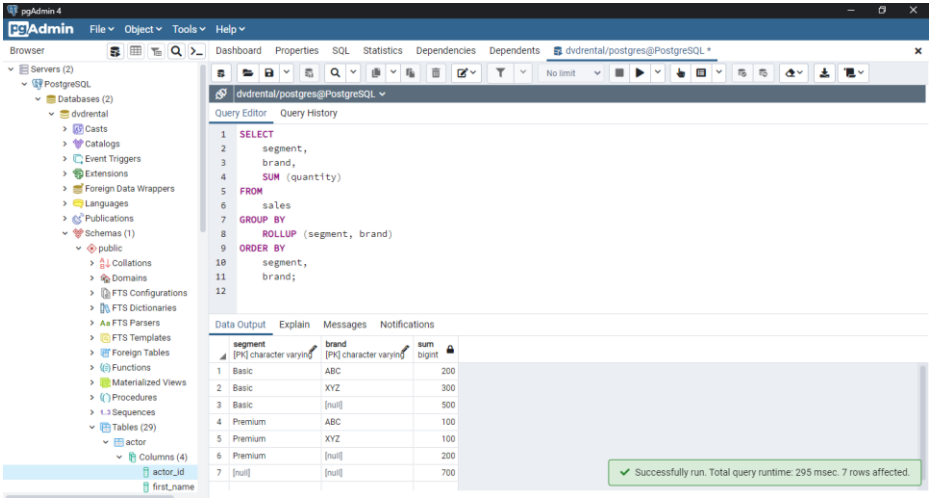


The following query uses the ROLLUP clause to find the number of products sold by brand (subtotal) and by all brands and segments (total).



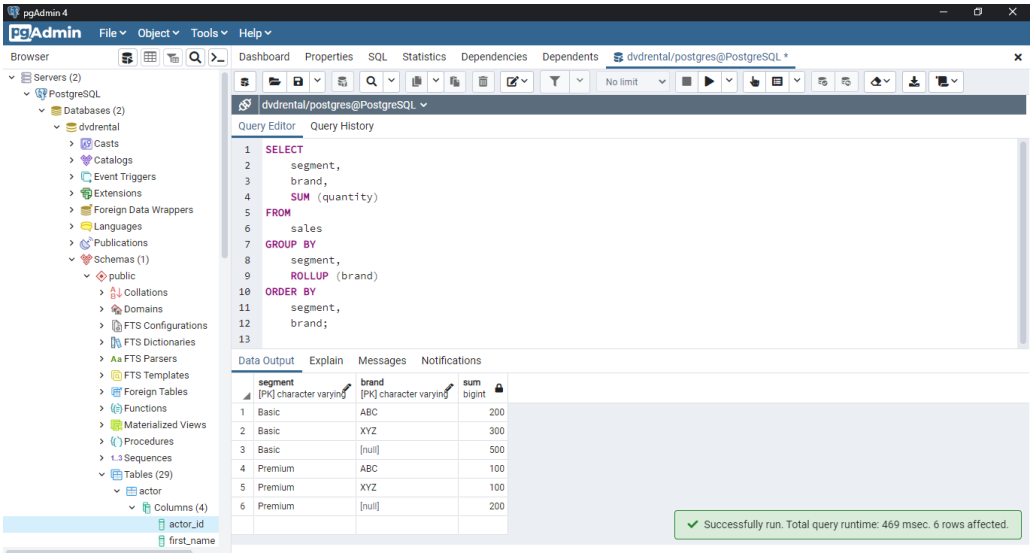
As you can see clearly from the output, the third row shows the sales of the ABC brand, the sixth row displays sales of the XYZ brand. The last row shows the grand total for all brands and segments. In this example, the hierarchy is brand > segment.

If you change the order of brand and segment, the result will be different as follows:



In this case, the hierarchy is the segment > brand.

The following statement performs a partial roll-up:



The following statement finds the number of rental per day, month, and year by using the ROLLUP:

